

Predictive Modeling on Mosquitos Carrying the West Nile Virus

Introduction/Problem:

West Nile Virus (“WNV”) first started appearing in Chicago in 2002. By 2004, the City of Chicago and the Chicago Department of Public Health (“CDPH”) had established a comprehensive surveillance and control program that is still in effect today.

The goal for our project is to build a model that can accurately predict the presence of West Nile Virus given weather and testing data. In doing so, the model can also help us derive insights on potential patterns that may exist in how the virus spreads. These findings will ultimately benefit the CPDH in optimizing its resources and future strategy on containing WNV.

Note that all of this data is publicly available and was provided as part of a Kaggle competition that was run in 2015 (<https://www.kaggle.com/c/predict-west-nile-virus>).

Data Wrangling:

We start the process by cleaning and wrangling the two datasets we’ll eventually use to train our model. These will be referred to as the “Weather” data and the “Mosquito Trap” data respectively.

Weather Data Preview:

	Station	Date	Tmax	Tmin	Tavg	Depart	DewPoint	WetBulb	Heat	Cool	...	CodeSum	Depth	Water1	SnowFall	PrecipTotal	StnPressure	SeaLevel
0	1	2007-05-01	83	50	67	14	51	56	0	2	...		0	NaN	0.0	0.00	29.10	29.82
1	2	2007-05-01	84	52	68	NaN	51	57	0	3	...		NaN	NaN	NaN	0.00	29.18	29.82
2	1	2007-05-02	59	42	51	-3	42	47	14	0	...	BR	0	NaN	0.0	0.00	29.38	30.09
3	2	2007-05-02	60	43	52	NaN	42	47	13	0	...	BR HZ	NaN	NaN	NaN	0.00	29.44	30.08
4	1	2007-05-03	66	46	56	2	40	48	9	0	...		0	NaN	0.0	0.00	29.39	30.12

Mosquito Trap Data Preview:

	Date	Address	Species	Block	Street	Trap	AddressNumberAndStreet	Latitude	Longitude	AddressAccuracy	NumMosquitos	WnvPresent
0	2007-05-29	4100 North Oak Park Avenue, Chicago, IL 60634,...	CULEX RESTUANS	41	N OAK PARK AVE	T002	4100 N OAK PARK AVE, Chicago, IL	41.954690	-87.800991	9	1	0
1	2007-05-29	4100 North Oak Park Avenue, Chicago, IL 60634,...	CULEX RESTUANS	41	N OAK PARK AVE	T002	4100 N OAK PARK AVE, Chicago, IL	41.954690	-87.800991	9	1	0
2	2007-05-29	6200 North Mandell Avenue, Chicago, IL 60646, USA	CULEX RESTUANS	62	N MANDELL AVE	T007	6200 N MANDELL AVE, Chicago, IL	41.994991	-87.769279	9	1	0
3	2007-05-29	7900 West Foster Avenue, Chicago, IL 60656, USA	CULEX RESTUANS	79	W FOSTER AVE	T015	7900 W FOSTER AVE, Chicago, IL	41.974089	-87.824812	8	1	0
4	2007-05-29	7900 West Foster Avenue, Chicago, IL 60656, USA	CULEX RESTUANS	79	W FOSTER AVE	T015	7900 W FOSTER AVE, Chicago, IL	41.974089	-87.824812	8	4	0

Data Wrangling Part 1: Weather Data

For the columns which denote precipitation ("Snowfall" & "PrecipTotal"), there are two values we need to replace. M denotes "Missing Data" which we should replace with NaN so that it does not interfere with our ability to perform numerical analysis on these columns. T denotes "Trace" which indicates that the value is greater than zero but less than the smallest unit of measurement (0.1 inches for Snowfall and 0.01 inches for rain). We will replace "T" with a value equal to half of the smallest unit (i.e. 0.05 inches for Snowfall and 0.005 inches for rain). At this time, we can also delete our "Water1" & "Depth" columns because these respective series are empty for our dataset.

Next, we convert all of the numerical features into Float datatypes. For the "Sunrise" and "Sunset" columns, we will convert these into datetimes. Then, with our weather conditions ("CodeSum"), we split these categories into indicator variables using Panda's `get_dummies()` function.

Lastly, we take the "Station" column and merge the "Station1" and "Station2" rows for each date by using the following methodology:

1. If there are any null values in one station dataset and not the other, then the merged version will use whatever is available.
2. All numerical values will be averaged.
3. For the weather condition categories, we will include every observed value between the two stations (e.g. If Station 1 recorded "BR" for and Station 2 recorded "BR HZ", then the merged row's "CodeSum" will have "BR HZ")

Our final weather dataset now has 35 columns.

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2944 entries, 0 to 2943
Data columns (total 22 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Station     2944 non-null  int64
1   Date        2944 non-null  object
2   Tmax        2944 non-null  int64
3   Tmin        2944 non-null  int64
4   Tavg        2944 non-null  object
5   Depart      2944 non-null  object
6   DewPoint    2944 non-null  int64
7   WetBulb     2944 non-null  object
8   Heat        2944 non-null  object
9   Cool        2944 non-null  object
10  Sunrise     2944 non-null  object
11  Sunset      2944 non-null  object
12  CodeSum     2944 non-null  object
13  Depth       2944 non-null  object
14  Water1      2944 non-null  object
15  SnowFall    2944 non-null  object
16  PrecipTotal  2944 non-null  object
17  StnPressure  2944 non-null  object
18  SeaLevel    2944 non-null  object
19  ResultSpeed  2944 non-null  float64
20  ResultDir    2944 non-null  int64
21  AvgSpeed    2944 non-null  object
dtypes: float64(1), int64(5), object(16)
memory usage: 506.1+ KB
```

Initial Weather Dataset

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1472 entries, 0 to 1471
Data columns (total 36 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        1472 non-null  object
1   Tmax        1472 non-null  float64
2   Tmin        1472 non-null  float64
3   Tavg        1472 non-null  float64
4   Depart      1472 non-null  float64
5   DewPoint    1472 non-null  float64
6   WetBulb     1469 non-null  float64
7   Heat        1472 non-null  float64
8   Cool        1472 non-null  float64
9   Sunrise     1472 non-null  object
10  Sunset      1472 non-null  object
11  CodeSum     1472 non-null  object
12  SnowFall    1472 non-null  float64
13  PrecipTotal  1472 non-null  float64
14  StnPressure  1470 non-null  float64
15  SeaLevel    1467 non-null  float64
16  ResultSpeed  1472 non-null  float64
17  ResultDir    1472 non-null  int64
18  AvgSpeed    1472 non-null  float64
19  CodeSumSplit 1472 non-null  object
20  BCFG        1472 non-null  int32
21  BR          1472 non-null  int32
22  DZ          1472 non-null  int32
23  FG          1472 non-null  int32
24  FG+         1472 non-null  int32
25  FU          1472 non-null  int32
26  GR          1472 non-null  int32
27  HZ          1472 non-null  int32
28  MIFG        1472 non-null  int32
29  RA          1472 non-null  int32
30  SN          1472 non-null  int32
31  SQ          1472 non-null  int32
32  TS          1472 non-null  int32
33  TSRA        1472 non-null  int32
34  VCFG        1472 non-null  int32
35  VCTS        1472 non-null  int32
dtypes: float64(14), int32(16), int64(1), object(5)
memory usage: 322.1+ KB
```

Final Weather Dataset

Data Wrangling Part 2: Mosquito Traps Dataset


First, we merge our cleaned weather dataframe into our mosquito trap dataframe using the “Date” column as our ID. We then delete any of the weather condition columns which sum to 0 because this indicates that for the dates we’re working with in our mosquito trap dataset, there were no instances of those weather conditions occurring.

Next, we use `get_dummies` again to separate the “Species” column into identifying variables like we did with the weather conditions.

Exploratory Data Analysis:

We group our dataframe by date and then sum the values which allows to create a ratio to track the trend of the virus over time.

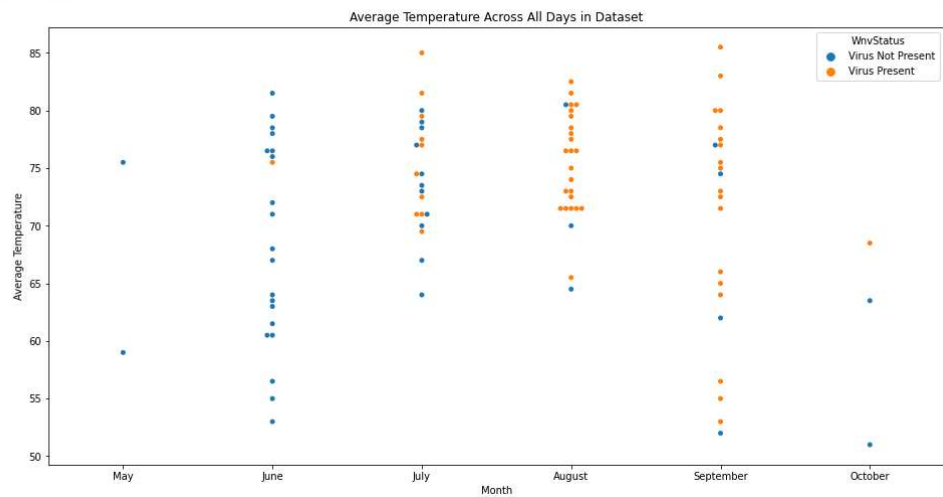
West Nile Virus Positive Tests Over Time as a Ratio to Total Traps & Total Mosquitos




Here, we can immediately see that the virus starts to pick up in July before peaking in August and then tapering off in September and October.

Using our same aggregated dataset, this time we take the mean across each date in order to observe trends between daily temperatures and the virus.

```
In [50]: df_mean = df_final.groupby('Date').mean()
df_mean['month'] = months
df_mean['year'] = years
df_mean['WnvStatus'] = virus
df_mean['NumMosquitos'] = df_mean['NumMosquitos'].astype('float')
swarmplot = sns.swarmplot(x='month', y='Tavg', hue='WnvStatus', data=df_mean)
plt.title('Average Temperature Across All Days in Dataset')
plt.ylabel('Average Temperature')
plt.xlabel('Month')
swarmplot.set_xticklabels(['May', 'June', 'July', 'August', 'September', 'October'])
plt.show();
```



Lastly, we make use of the geopandas library to represent our data on a map of Chicago where we can see that Northwest Chicago appears to be a hotspot for the virus.



Feature Engineering:

We add the following features to our dataframe:

1. Days of the Week
2. Month & Year
3. Days Since Previous Weather Condition: For each weather condition, we mark the number of days that pass between two consecutive instances of that weather condition occurring. Note that because there are two year gaps in our data (2007, 2009, 2011, 2013), we also have to take care to mark off the start of a new year with “N/A”.
4. Municipalities: Using geocoders, we input our latitude and longitude values into Nominatim to pull up details on each of our trap locations. Municipalities was selected as the new categorical feature to be added because it had fewer null entries than other potential features while also having less variety in its actual values which helps prevent our data’s dimensionality from growing exponentially large.

All of these added features are categorical variables so we run through `get_dummies` again to convert them all into indicators. After this step, we address all of the missing data by imputing any null values with the mean of their respective series.

```
In [14]: from sklearn.impute import SimpleImputer
df_final = df2.copy(deep=True)
mean_imputer = SimpleImputer(strategy='mean')
df_final.loc[:,['WetBulb','StnPressure', 'TimeSinceLastBR',
               'TimeSinceLastDZ','TimeSinceLastFG',
               'TimeSinceLastHZ','TimeSinceLastRA','TimeSinceLastTS',
               'TimeSinceLastTSRA','TimeSinceLastVCTS']] = mean_imputer.fit_transform(
    df2[['WetBulb','StnPressure', 'TimeSinceLastBR',
         'TimeSinceLastDZ','TimeSinceLastFG',
         'TimeSinceLastHZ','TimeSinceLastRA','TimeSinceLastTS',
         'TimeSinceLastTSRA','TimeSinceLastVCTS']])
```

Model Preprocessing:

Now we can split our data into training and test sets. However, because there are an overwhelming number of “negative” WNV cases in our data, we need to sample an equal proportion of both positive and negative cases so that our model can fairly learn both (otherwise we end up with a model that will inevitably predict “negative” on almost everything). After taking this sample, we feed it into sklearn’s `train_test_split`.


Our last step before entering the modeling stage is to reduce the dimensionality of our data by only selecting the features that will have the most impact in our model predictions. We do this by calculating the information value of each feature and then only selecting the ones which have $IV > 0.01$ (has some amount of predictive power) but also < 0.80 (is suspiciously high and would overrule the other features). We also calculate the variance inflation factor (“VIF”) on each of these features and remove any which have a value > 5 as this would be an indication that the feature in question has high collinearity with another feature; this process helps keep the features independent.

At the end of our preprocessing step, we are left with 18 features in our final dataset.

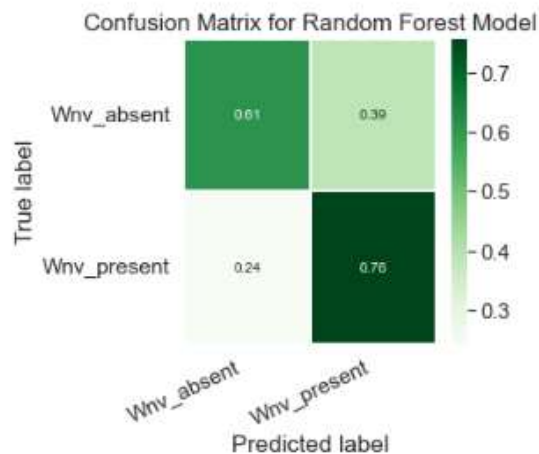
Modeling & Results:

We built a Random Forest Model and an XG Boost Model. The performance metrics for each are presented below:


Random Forest:




Random Forest: Accuracy=0.680
Random Forest: F1 Score=0.679
Random Forest: ROC Score=0.685





XG Boost:



XG Boost: Accuracy=0.695
XG Boost: f1-score=0.695
XG Boost: ROC Score=0.699



Our XG Boost model does slightly better overall on accuracy, f1-score, and ROC score so it is selected as the final model for our project. However, note that Random Forest actually slightly outperformed XG Boost on recall. Future iterations of this project will likely explore new methodologies that can improve the random forest model on the basis that true positives are the most important metric to optimize given the context of this dataset.



Our SHAP analysis reveals that the most impactful features in our model were as follows:

1. August: If the data was from August, this made the model more likely to predict positive.
2. TimeSinceLastTSRA: Generally, the longer it has been since there was last a rainy thunderstorm, the more likely the model predicts positive.
3. Departure from the 30-year normal temperature: The more extreme the temperature, the more likely the model predicts positive. This is potentially related to why August is the peak month since that's when the highest temperatures occur.
4. Jefferson Township: Traps from Jefferson Township (which encompasses most of Northwest Chicago) will make the model more likely to predict positive.

Note that these features corroborate with our earlier observations made during our exploratory data analysis.

Final Thoughts & Recommendations:

Although the CPDH has been monitoring the WNV situation since 2004, the prevalence of the virus has only continued to trend upwards year by year. Our model results produce the following recommendations on strategies to optimize the city's efforts in combating this virus for future years:

1. Increase coverage of spraying in the Northwest Chicago region particularly the areas encompassed in the Jefferson Township Municipality.
2. Focus the spraying schedule to be more concentrated in August.
3. Add in additional out-of-schedule sprays if temperatures are noted to be particularly high for a certain day OR if at least 3 weeks have gone by without any rain or thunderstorms.

Our final model had an overall ROC score of 0.699 and a recall score of 0.75. Future efforts to improve our model will focus on trying to improve the recall since the implications of positive cases are far greater than the consequences of misclassifying negative cases as false positives.